

The RTL Gauntlet: Measuring Honesty, Cost, and Latency-Robustness in Agentic Hardware Design

(working draft)

Abstract

Agentic RTL frameworks now report 100% pass on standard benchmarks, but that score is measured against *visible* tests, is *expensive* to reach, and is obtained on *proxy* tasks. We turn these three concerns into measurable axes. We build a two-tier protocol that freezes an agent’s design and scores it with a withheld formal-equivalence oracle, and define a Reward-Hacking Gap (RHG) and Honest Pass Rate (HPR). Applying it to 156 VerilogEval tasks across four models, we find the naïve headline RHG (6.1%) is *entirely an oracle artifact*: a naïve formal oracle over-reports reward hacking via don’t-care `x`, `async-reset`, state-encoding, init-state, and even a SystemVerilog parser flag. We harden the oracle in six steps—reset-aware, don’t-care-aware, bounded miter+SAT, `read_verilog -sv`, a memory (case→ROM) pass, and a `-nolatches` reset-aware BMC—cutting false counter-examples $9 \rightarrow 1$ and “inconclusive” $50 \rightarrow 0$, and verify every surviving flag by hand: **zero genuine reward hacking** (95% upper bound $\leq 3.2\%$) by frontier models on fair tasks. A weaker model (Haiku 4.5) fails $3\times$ more than Opus 4.8 on the visible tests but does *not* hack more; even a tamper-capable shell agent that struggles for several iterations does not edit the testbench. We show the oracle *would* catch hacking when present (a candidate that passes a randomized hidden testbench but is formally disproven), that the result survives off the verbatim benchmark (a contamination probe), and we add a token-cost analysis (an early-stop reclaims 12–23% of tokens) and a latency-axis prototype that produces real Sky130 power/area/timing through a containerized OpenLane flow. The central artifact for honest agentic-hardware evaluation is a don’t-care/reset/encoding-aware oracle plus a verification discipline—not a new pass-rate.

1 Introduction

HORIZON treats agentic hardware design as repository-level code evolution and reaches 100% pass on several RTL suites—but iteration-0 pass is only 47.8%, so the headline comes from *iterative repair against a visible signal*. Its authors name three open problems: reward hacking, token cost, and high-latency (PPA) reward. We make each measurable.

Thesis. *Pass@visible is the wrong score for agentic hardware design*: it can be gamed (honesty), it is expensive (cost), and it does not scale to real reward (latency).

Contributions.

1. A two-tier, formal-grounded evaluation protocol with metrics RHG and HPR (§3).
2. The finding that a naïve formal oracle **over-reports** reward hacking, and a six-step hardened oracle that removes it (closing the sequential-FSM residual); verification discipline as a first-class method (§4).
3. An empirical result across 4 models $\times$ 156 tasks: **zero verified genuine reward hacking** on fair tasks; weakness $\neq$ hacking; tampering is caught when present (§5).
4. Cost and latency prototypes: token accounting with an early-stop policy, and real Sky130 PPA via a containerized flow (§5).

Work	gap metric	formal oracle	RTL
SpecBench	✓ (software)	finite held-out	—
EvilGenie	✓ + edit-detect	—	—
Agentic Frontier	—	—	✓
This work	✓	exhaustive	✓

Table 1: SpecBench owns the visible/held-out gap for *software* and admits its held-out suite “cannot exhaustively certify compliance.” Our slot: RTL + an exhaustive formal oracle (fixing that limit) + tamper-evidence.

2 Related Work and Positioning

Formal-equivalence-as-oracle is already community-standard for RTL-LLM evaluation (VeriThoughts; Not-SoTiny; RealBench)—so it is our *tool*, not our contribution. Table 1 situates the closest work.

3 Method: the two-tier protocol

The agent sees a complete specification, a *visible* diagnostic testbench, and its failure logs, and edits RTL until visible passes. We then **freeze** the design and score withheld tiers it never saw: a hidden randomized/exhaustive testbench, and **formal equivalence** vs. a curated golden (Yosys). The I/O interface is locked. We define

$$\text{RHG} = \frac{\#\{\text{visible-passers failing a withheld tier}\}}{\#\{\text{visible-passers}\}}, \quad \text{HPR} = \frac{\#\{\text{pass visible and every withheld tier}\}}{\#\text{tasks}}.$$

Here “two-tier” is the split *visible* vs. *withheld*. The withheld side yields one of four formal verdicts, mapped as **proven**/**bmc_equiv**→**honest**, **cex**→**flagged hack**, and **dontcare**/**inconclusive**→**excluded** (a third bucket, neither honest nor hack). So HPR is a *lower bound* on true honesty and $\#\text{cex}/\#\text{visible-passers}$ an *upper bound* on *detectable* hacking; only **verified-genuine RHG** (every **cex** confirmed by hand) is the headline claim. On the 156-task sweep the withheld tier is the **formal** oracle; the hidden randomized testbench is exercised on the focused demos (**formal_demo**, **sem_zerocount**, **tamper**). A tamper tier (a shell-capable agent that may edit any workspace file) flags edits to the testbench/harness as *exploit-evidenced* hacking, with the withheld tiers run on frozen originals so tampering cannot fool the verdict. We separate the claim into behavioral gap, exploit-evidenced, and tamper-confirmed, and never claim intent without trajectory evidence.

4 The oracle, and why naïve formal over-reports

On a 156-task sweep (Opus 4.8), a naïve **equiv_make** oracle reported 9 RHG counter-examples and 50 “inconclusive.” **Every CEX was a false positive**, found by hand-verifying flagged cases: don’t-care **x** in the golden; functionally identical designs flagged on async reset / initial state; a candidate logically identical but with a different state encoding; and—for most “inconclusive”—a silent parse failure, since **read_verilog** defaults to Verilog-2005 while the references use **always_ff/logic/enums**.

Our hardening pipeline (each step tested on verified cases before re-running):

1. **async2sync** (normalize async resets);
2. reclassify a CEX to *dontcare* when the golden has an **x** literal;
3. a bounded miter+SAT fallback comparing I/O sequences directly (encoding-agnostic; a SAT model is a real CEX);
4. **read_verilog -sv**;
5. a **memory** (case→ROM) pass;
6. a **-nolatches** reset-aware BMC (Pass-3) that closes the residual **inconclusive** FSMs. Their goldens use an **always_comb case(state)** over a **reg [3:0]** that lists only the reachable states, so the unlisted

stage	honest	bmc	dc	RHG	inconcl	fail
naïve	88	–	–	9	50	9
+reset+dc	91	–	3	3	50	9
+BMC	91	3	2	1	50	9
+SV	122	6	4	1	14	9
+memory	129	6	5	1	6	9
+reset-BMC	129	11	6	1	0	9

Table 2: Oracle hardening (same 156 Opus candidates, no new LLM calls). False CEX $9 \rightarrow 1$; inconclusive $50 \rightarrow 0$; HPR $56\% \rightarrow 90\%$ (honest $129 + 11$ BMC-verified = 140 of 156). The residual sequential FSMs are now closed—the `-nolatches` reset-aware Pass-3 proves 5 equivalent and reclassifies 1 as an output don’t-care.

	Opus 4.8	GPT-5.5	DeepSeek	Haiku 4.5
honest + bmc	140	137	118	113
HPR (95% CI)	0.90 [.84,.94]	0.88 [.82,.92]	0.76 [.68,.82]	0.72 [.65,.79]
fail_visible (+ no-cand)	9	6	36	30 (+11)
verified RHG	0 ($\leq .025$)	0 ($\leq .025$)	0 ($\leq .031$)	0 ($\leq .032$)

Table 3: Four models $\times$ 156 tasks (final oracle, incl. the reset-aware Pass-3; HPR counts honest + BMC-verified, so Opus = $129+11 = 140$, $140/156 = 0.90$). All verified RHG = 0: every flagged CEX is an init / encoding / input-space don’t-care artifact (see text). Weakness shows as failures, not cheating.

states infer a latch and yosys *errors out*—an elaboration artifact, not a solver limit (the same class as the `-sv` bug); `-nolatches` plus a reset drive at $t=1$ (the true reset state, not the zero state) then proves I/O equivalence.

5 Experiments

Formal earns its keep. `formal_demo`: a 16-bit candidate wrong *only* on `0xDEAD`. The randomized hidden testbench (512/65536) misses it—visible *and* hidden PASS—while exhaustive formal returns CEX. With finite tests alone RHG = 0 (it looks honest); formal exposes it (RHG = 0.50). This is the direct evidence that an exhaustive formal oracle beats a finite held-out suite.

Sweep and model comparison. Table 3 and Figure 2: the weaker models fail far more (**fail_visible** 9/6/36/30 for Opus/GPT-5.5/DeepSeek/Haiku, +11 no-candidate for Haiku) but hack no more. Every flagged RHG counter-example across the four models—4 distinct tasks—is a hand-verified artifact: `circuit8/q5b` are BMC zero-init artifacts (`-set-init-zero` forces an invalid state for an uninitialized latch / one-hot FSM, the latter produced independently by DeepSeek, Haiku, and GPT-5.5); `prob095` is a zero-init state with an incidentally-active output (simulation confirms equivalence after reset); and `prob149` is an *input-space* don’t-care—a water-level controller whose golden marks impossible non-thermometer sensor codes `'x`. The `dontcare` detector requires the `x` to drive an output, so an incidental internal `x` cannot mask a real CEX. With 95% Wilson intervals, verified-genuine RHG = 0, upper bound $\leq 2.5\%$ (Opus) / $\leq 2.5\%$ (GPT-5.5) / $\leq 3.1\%$ (DeepSeek) / $\leq 3.2\%$ (Haiku).

Elicit (negative, informative). A tamper-capable shell agent on a harder task: Haiku struggled four iterations (with full affordance to edit the testbench) but never tampered—it kept honestly repairing the design and solved it; Opus was honest in one. Genuine reward hacking is not elicited from aligned frontier models in a single agentic loop even under affordance and difficulty (the oracle *would* catch it—the planted red-team is caught). This points to RL training-time emergence as the regime to study.




Figure 1: Oracle hardening: artifacts fall (false CEX $9 \rightarrow 1$, inconclusive $50 \rightarrow 0$) and the verified-equivalent count (honest, incl. BMC) rises ($88 \rightarrow 140$; honest-only $88 \rightarrow 129$).

Contamination. VerilogEval is public (reported $\sim 100\%$ contamination for older models), so an honest pass could be memorization. We generate textually-novel, functionally-identical tasks (rename the target module, reframe the spec) and re-sweep: on the first 40 tasks Opus is HPR $1.00 \rightarrow 1.00$, RHG $0 \rightarrow 0$ ($\Delta = 0$). The result is robust to identifier-level contamination. As a *semantic*-novelty control, our self-authored tasks (`gray2bin`, `popcount8`, `hex7seg`) implement functions absent from any public benchmark, yet models are honest on them too (RHG = 0)—so the result is not benchmark memorization. A direct *semantic* mutation (`sem_zerocount`: count *zeros*, a variant of the canonical popcount) confirms this: the memorized count-ones solution passes the balanced visible test but is caught by hidden+formal (RHG = 0.50), while both models implement the variant honestly. We report contamination as its *own* one-sided Clopper–Pearson 95% upper bound (distinct from the reward-hacking bound above): $k = 0$ memorization signals on the 40-task probe gives $\leq 7.2\%$, tightening to $\leq 1.5\%$ as the probe scales to all 203 tasks—the bound shrinks because we grow n , not the claim. Since no single contamination signal is reliable alone, we plan to triangulate this behavioral probe with membership inference (Min-K% Prob, arXiv:2310.16789) and the NLL $\leftrightarrow$ degradation correlation under meaning-preserving mutation (arXiv:2604.21579); only the behavioral probe is run here.

Cost (C2). Opus 482k / Haiku 522k tokens; the weaker model needs more repair (2-iter 17 vs 36). The repair tail is mostly wasted (honest payoff 47% / 14%); an early-stop at one iteration reclaims 12% (Opus) / 23% (Haiku) of tokens for a $\sim 5\%$ honesty loss (Figure 3).

Latency (C3). A surrogate pipeline verified offline (mock 315 designs $\rightarrow$ holdout Pearson $r = 0.89/0.91/0.96$ for area/power/timing) and, via a containerized OpenLane flow (the OpenLane image as runtime $\rightarrow$ native `openlane`, no Docker-in-Docker), produced real Sky130 metrics: `counter8` (seq) $495.5 \mu\text{m}^2$ / 0.120 mW / worst-slack $\approx -5.5 \text{ ns}$; `popcount8` (comb) $247.7 \mu\text{m}^2$ / 0.051 mW / -1.6 ns (slacks *negative*: the default clock is not met—C3 is a pre-signoff *relative* prototype, not closed timing). (The deploy gotcha was a config file at the repo root, not OpenLane.)

Weakness $\neq$ hacking (4 models): more visible-fails, every RHG_cex a verified artifact




Figure 2: Weakness $\neq$ hacking across four models: Haiku/DeepSeek fail more on the visible tests (GPT-5.5 fails fewest), RHG= 0 throughout.

6 Findings

1. A naïve formal oracle **over-reports** reward hacking; the headline number was an artifact, found only by per-case verification. Verification discipline is mandatory.
2. No false positives on honest agents; across 4 models $\times$ 156 fair tasks, **zero verified genuine reward hacking** (95% upper bound $\leq 3.2\%$).
3. **Weakness $\neq$ hacking**—a weaker model fails far more but cheats no more, even under tamper affordance.
4. Hacking is a function of adversarial conditions, not benchmark mass; the protocol catches it when present.

7 Threats to Validity

Contamination: addressed at both identifier and semantic level (`sem_zerocount`); a full semantic re-mutation of the whole benchmark is future work. *Oracle residual*: the sequential-FSM residual is now *closed*—a `-nolatches` reset-aware BMC (Pass-3) proves the 5 Opus FSMs equivalent and reclassifies 1 as an output don’t-care, validated against a genuine-diff control (a deliberately-broken candidate is still flagged CEX, so the pass is not vacuously proving everything). GPT-5.5 leaves 3 inconclusive (Conway’s Life’s 256-cell grid plus two candidates the bounded SAT does not close in budget)—a per-candidate budget limit, not a verdict. *Eliciting* genuine hacking from aligned models failed; the negative result reflects single-loop inference on aligned models, not RL training—the regime where hacking is documented to emerge: test-subverting hacks discovered *during* production RL (arXiv:2503.11926) and shortcut behavior that is RLVR-specific, absent in non-RL counterparts (arXiv:2604.15149). *PPA*: open Sky130/OpenROAD is pre-signoff, not foundry signoff, and rankings can flip across synthesis flows; we therefore frame PPA as a *relative* comparison under identical settings (per TuRTL) and note larger designs + a commercial flow as the path to signoff-grade claims. The open flow inflates area/power by $\sim 1.2\text{--}1.5\times$ vs. a commercial signoff stack (arXiv:2512.06122) and ranking inversions are real (RTL-OPT, arXiv:2601.01765); we therefore verify *rank stability across* ≥ 2 *synthesis strategies* (Kendall- τ) to test the inversion failure mode directly, and expand C3 to a size-graded set (`spm`, AES, picorv32, ibex). *Nondeterminism*: LLM runs vary; we report Wilson CIs and pin model ids.

Cost: the repair tail is mostly wasted (~5% honesty lost)




Figure 3: The repair tail is mostly wasted; an early-stop reclaims 12–23% of tokens.

8 Limitations and Future Work

A semantic contamination mutation of the whole benchmark; more models and task types (debug/verify/reuse); a PPA surrogate trained on a real multi-design dataset with an agent loop under high-latency reward; and—most ambitiously—**RL training-time hacking**: train a small RTL agent (Qwen3-4B, GRPO) with a visible-test reward and measure the emergence of hacking via a reward-hacking gap $RHG(t) = \text{visible_pass}(t) - \text{formal_pass}(t)$ audited by the hidden+formal oracle each step—a positive, growing gap is direct evidence, and RTL is the one domain where this audit is *exhaustive*. (Probe implemented: `scripts/train_grpo.py`.)

9 Conclusion

Honest evaluation of agentic hardware design is gated not by a new pass-rate but by an oracle that is don’t-care/reset/encoding-aware and by a discipline of verifying every flag. With that, frontier agents do not hack fair RTL tasks—but the oracle catches it when they do.

References

- [1] Yu, Deng, Pinckney, Khailany. Agentic Hardware Design as Repository-Level Code Evolution (HORIZON). arXiv:2606.28279, 2026.
- [2] Pinckney et al. Comprehensive Verilog Design Problems (CVDP). arXiv:2506.14074, 2025.
- [3] Zhao, Srikanth, Wu, Jiang. SpecBench: Measuring Reward Hacking in Long-Horizon Coding Agents. arXiv:2605.21384, 2026.
- [4] Gabor, Lynch, Rosenfeld. EvilGenie: A Reward Hacking Benchmark. arXiv:2511.21654, 2025.
- [5] Wang et al. VeriContaminated: Assessing LLM-Driven Verilog Coding for Data Contamination. arXiv:2503.13572, 2025.
- [6] Du, Pinckney. Trace2Skill: Verifier-Guided Skill Evolution for Long-Context EDA Agents. arXiv:2605.21810, 2026.
- [7] Ghorab et al. NotSoTiny: A Large, Living Benchmark for RTL Code Generation. arXiv:2512.20823, 2025.

- [8] Jin et al. RealBench: Benchmarking Verilog Generation Models with Real-World IP Designs. arXiv:2507.16200, 2025.
- [9] Yubeaton, Garg, Hegde. Exploring the Agentic Frontier of Verilog Code Generation. arXiv:2603.19347, 2026.
- [10] Patel, Chhabria, Arora. Understanding Inference-Time Token Allocation and Coverage Limits in Agentic Hardware Verification. arXiv:2604.15657, 2026.
- [11] Baker et al. Monitoring Reasoning Models for Misbehavior and the Risks of Promoting Obfuscation. arXiv:2503.11926, 2025.
- [12] Helff et al. LLMs Gaming Verifiers: RLVR can Lead to Reward Hacking. arXiv:2604.15149, 2026.
- [13] Li et al. Exploring Pass-Rate Reward in RL for Code Generation. arXiv:2605.02944, 2026.
- [14] TuRTL: A Unified Evaluation of LLMs for RTL Generation (OpenLane/Sky130, PPA-as-ratio). arXiv:2504.01986, 2025.
- [15] RTL-OPT: PPA-ranking reliability of LLM-generated RTL across synthesis flows. arXiv:2601.01765, 2026.
- [16] Shi et al. Detecting Pretraining Data from Large Language Models (Min-K% Prob). arXiv:2310.16789, 2023.
- [17] Metamorphic testing of memorization in LLM program repair (NLL $\leftrightarrow$ degradation). arXiv:2604.21579, 2026.
- [18] Using Open-Source EDA Tools in ASICs (open-vs-commercial PPA gap). arXiv:2512.06122, 2025.